



CS 584 PyTorch Tutorial 2

Jason Choi

https://github.com/boslbi92/pytorch_tutorial

Agenda

1. Neural Collaborative Filtering (NCF)
 - Movie Ratings Prediction
2. Implementing NCF with PyTorch
3. Coding tutorial

Movie Ratings Prediction

1. MovieLens dataset
 - [user_id, item_id, rating]
 - How to handle sparse input?
2. Neural collaborative filtering
 - $f(V_{user}, V_{item}) \rightarrow rating$

	i_1	i_2	i_3	i_4	i_5
u_1	1	1	1	0	1
u_2	0	1	1	0	0
u_3	0	1	1	1	0
u_4	1	0	1	1	1

↑ users

← items →

(a) user–item matrix

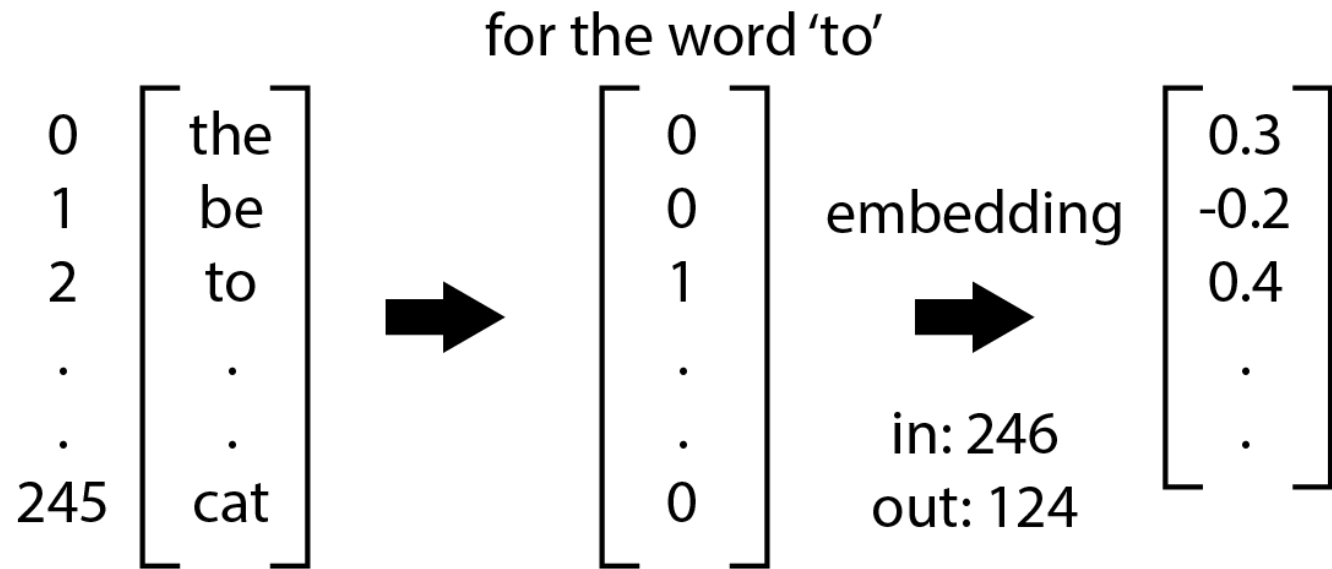
Neural Collaborative Filtering

1. Multi-layer perceptron (MLP)
 - Model interactions by non-linear transformations
2. Generalized matrix factorization (GMF)
 - Model interactions by learning latent user and item vectors
3. Neural collaborative filtering (NCF)
 - Concatenate outputs from first & second modules

Multi-layer Perceptron

- Represent user $\rightarrow V_{user}$
- Represent item $\rightarrow V_{item}$
- Represent interaction $\rightarrow [V_{user}; V_{item}]$
- Apply nonlinear transformations
 - N-layers of nonlinear transformations to model interaction

Use of Embeddings Layer



Assign random vectors to each user and item

- Represent user $\rightarrow V_{user}$
- Represent item $\rightarrow V_{item}$
- Represent interaction $\rightarrow [V_{user}; V_{item}]$

Multi-layer Perceptron

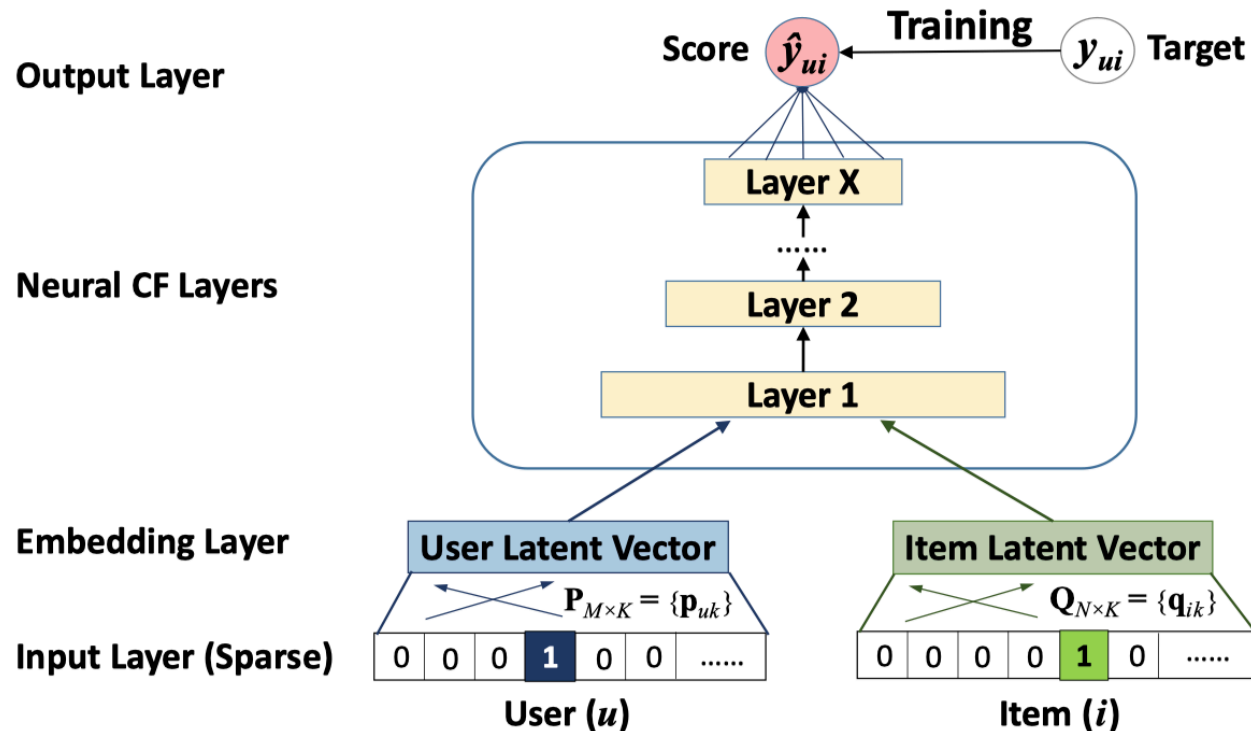


Figure 2: Neural collaborative filtering framework

1. User embedding layers
 - (num_user, latent_dim)
 - User $x \rightarrow x^{th}$ row of embedding matrix
 2. Item embedding layers
 - (num_item, latent_dim)
 - Item $y \rightarrow y^{th}$ row of embedding matrix
 3. 3-layered MLP
 4. Predict scalar value
- How to improve?

Generalized Matrix Factorization

- Learn good p_u and q_i embeddings to explain interactions

$$\hat{y}_{ui} = f(u, i | \mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u^T \mathbf{q}_i = \sum_{k=1}^K p_{uk} q_{ik},$$

p_u = latent vector for a user

q_i = latent vector for an item

\hat{y}_{ui} = inner products of p, q

Neural Collaborative Filtering

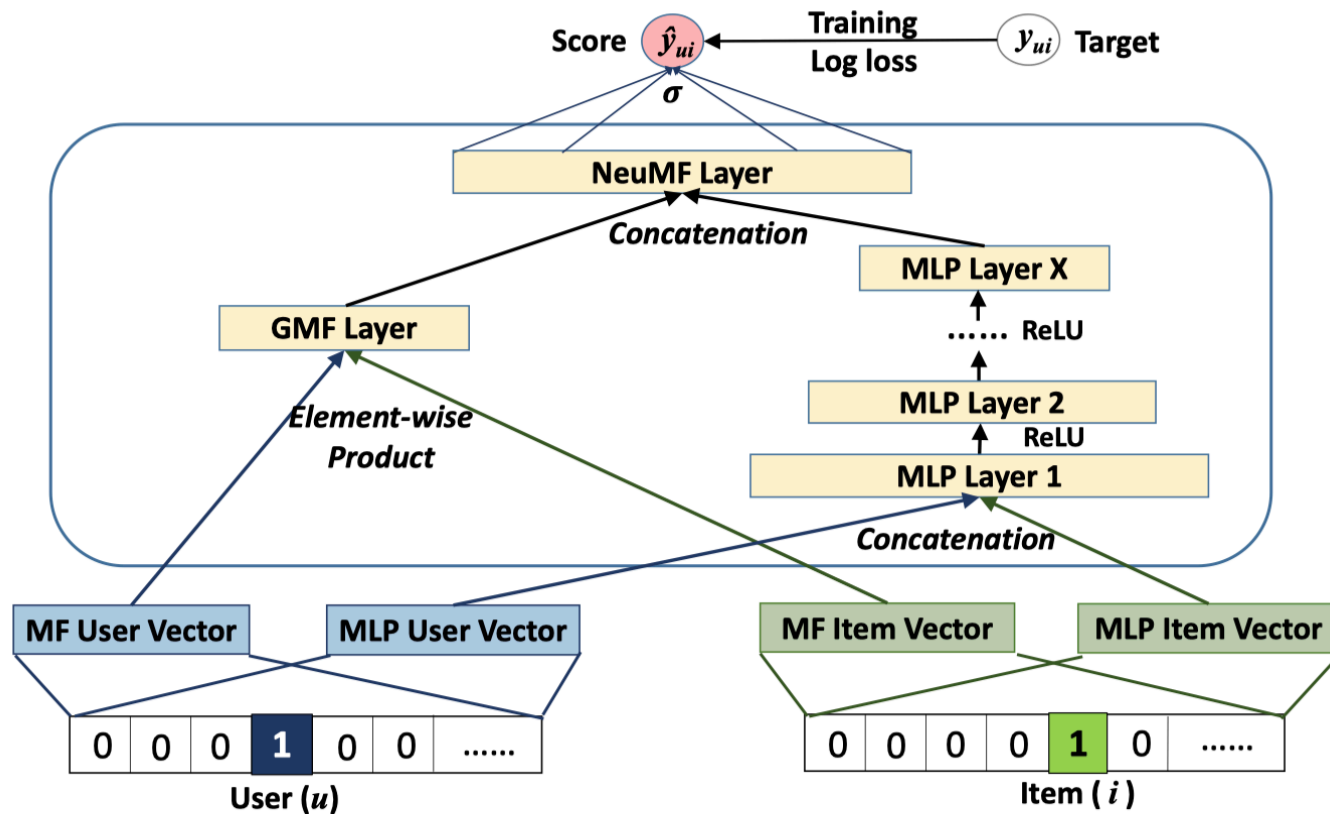



Figure 3: Neural matrix factorization model

1. GMF module
2. MLP module
3. Combine two modules
4. Predict scalar value

Agenda

1. Neural Collaborative Filtering (NCF)
 - Movie Ratings Prediction
2. Implementing NCF with PyTorch
3. Coding tutorial




Step 1: Prepare Inputs

- Convert dataset into tensor
- Initialize DataLoader class

```
# convert input to torch tensors
train_user = torch.tensor(train_data['user'].values,
                           device=device)
train_item = torch.tensor(train_data['item'].values,
                           device=device)
train_rating = torch.tensor(train_data['rating'].values,
                             device=device,
                             dtype=torch.float)

# convert tensors to dataloader
train_dataset = data.TensorDataset(train_user,
                                    train_item,
                                    train_rating)
train_loader = data.DataLoader(train_dataset,
                               batch_size=batch_size,
                               shuffle=True)
```



Step 2: MLP Module

- Initialize user embeddings
- Initialize item embeddings

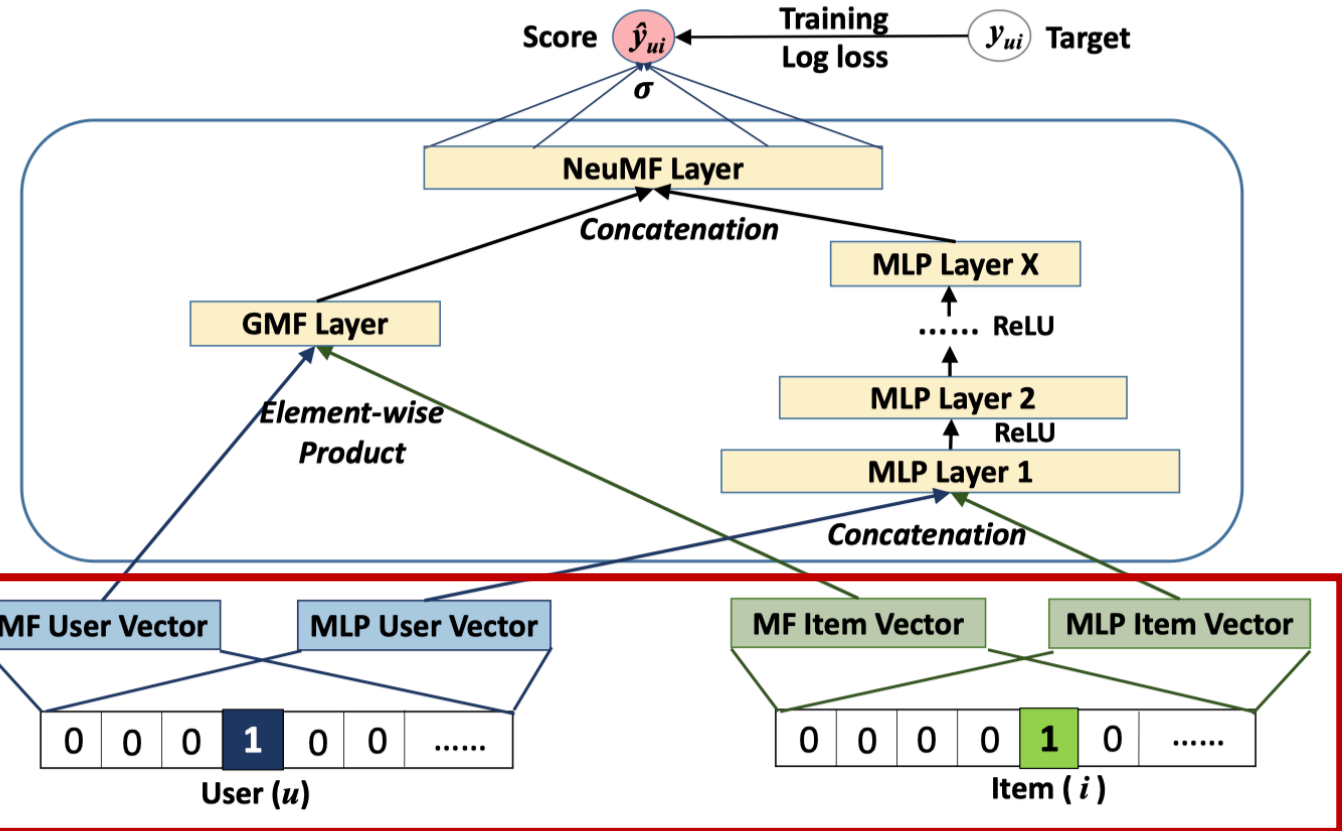



Figure 3: Neural matrix factorization model

```
# multilayer perceptron sub-module
class MLP(nn.Module):
    def __init__(self, user_num, item_num, output_dim):
        super(MLP, self).__init__()
        self.user_emb = nn.Embedding(user_num, output_dim * 4)
        self.item_emb = nn.Embedding(item_num, output_dim * 4)
```




Step 2: MLP Module

- Initialize linear layers
- 1st layer → 128
- 2nd layer → 64
- 3rd layer → 32

```
# multilayer perceptron sub-module
class MLP(nn.Module):
    def __init__(self, user_num, item_num, output_dim):
        super(MLP, self).__init__()
        self.user_emb = nn.Embedding(user_num, output_dim * 4)
        self.item_emb = nn.Embedding(item_num, output_dim * 4)

        self.linear_1 = nn.Linear(output_dim * 8, output_dim * 4)
        self.linear_2 = nn.Linear(output_dim * 4, output_dim * 2)
        self.linear_3 = nn.Linear(output_dim * 2, output_dim)
        self.dropout = nn.Dropout()
        self.relu = nn.ReLU()
```



Step 2: MLP Module

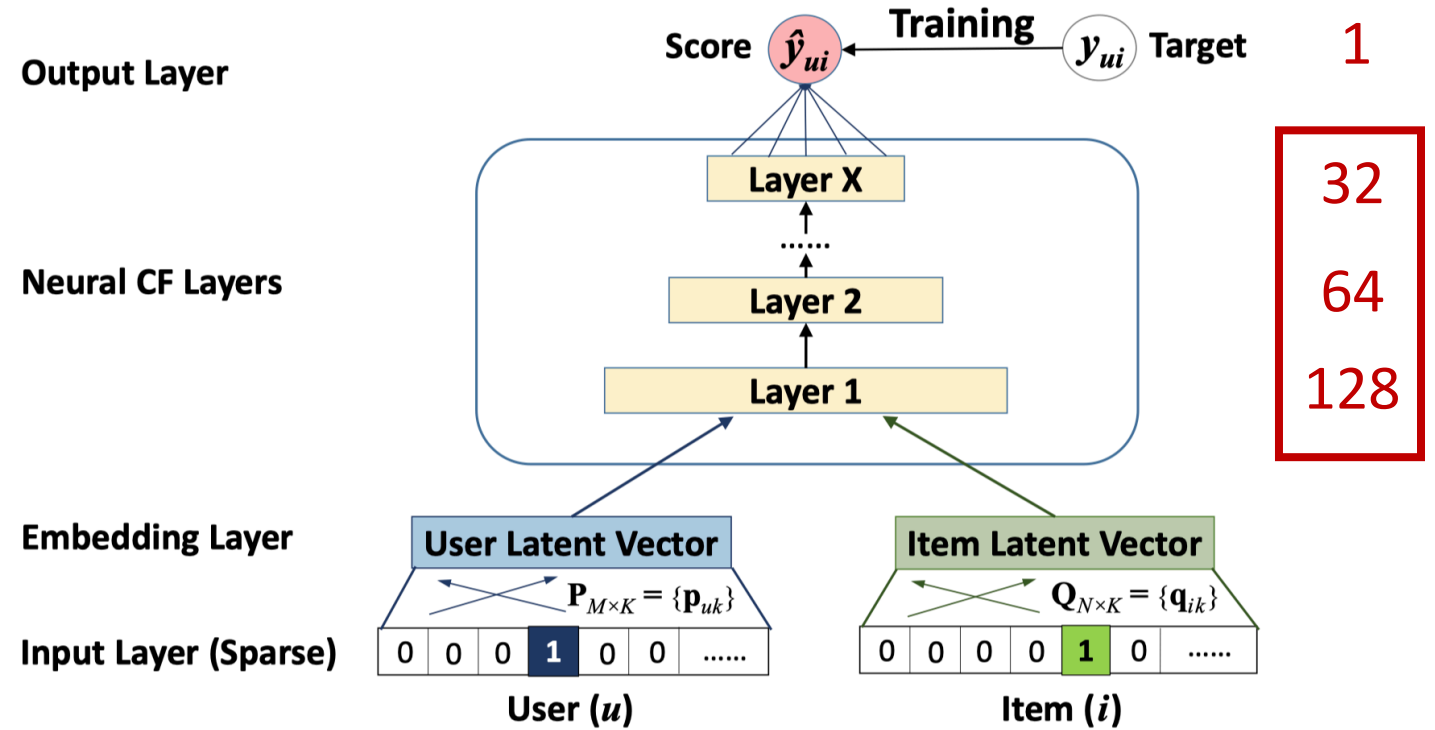




Figure 2: Neural collaborative filtering framework



Step 2: MLP Module

- Represent user-item interaction
- Three matrix multiplications

```
def forward(self, user, item):  
    user_emb = self.user_emb(user)  
    item_emb = self.item_emb(item)  
    concat = torch.cat((user_emb, item_emb), -1)  
  
    x = self.linear_1(concat)  
    x = self.relu(x)  
    x = self.dropout(x)  
  
    x = self.linear_2(x)  
    x = self.relu(x)  
    x = self.dropout(x)  
  
    x = self.linear_3(x)  
    x = self.relu(x)  
    return x
```




Step 3: GMF Module

- Initialize user embeddings
- Initialize item embeddings

```
# generalized matrix factorization sub-module
class GMF(nn.Module):
    def __init__(self, user_num, item_num, output_dim):
        super(GMF, self).__init__()
        self.user_emb = nn.Embedding(user_num, output_dim)
        self.item_emb = nn.Embedding(item_num, output_dim)
        self._init_weight()

    def forward(self, user, item):
        user_emb = self.user_emb(user)
        item_emb = self.item_emb(item)
        output = user_emb * item_emb
        return output
```

$$\hat{y}_{ui} = f(u, i | \mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u^T \mathbf{q}_i$$




Step 3: GMF Module

- Initialize user embeddings
- Initialize item embeddings

```
# generalized matrix factorization sub-module
class GMF(nn.Module):
    def __init__(self, user_num, item_num, output_dim):
        super(GMF, self).init_()
        self.user_emb = nn.Embedding(user_num, output_dim)
        self.item_emb = nn.Embedding(item_num, output_dim)
        self._init_weight_()

    def forward(self, user, item):
        user_emb = self.user_emb(user)
        item_emb = self.item_emb(item)
        output = user_emb * item_emb
        return output
```

$$\hat{y}_{ui} = f(u, i | \mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u^T \mathbf{q}_i$$




Step 3: GMF Module

- Initialize user embeddings
- Initialize item embeddings
- Output is product of two latent vectors

```
# generalized matrix factorization sub-module
class GMF(nn.Module):
    def __init__(self, user_num, item_num, output_dim):
        super(GMF, self).__init__()
        self.user_emb = nn.Embedding(user_num, output_dim)
        self.item_emb = nn.Embedding(item_num, output_dim)
        self._init_weight()

    def forward(self, user, item):
        user_emb = self.user_emb(user)
        item_emb = self.item_emb(item)
        output = user_emb * item_emb
        return output
```

$$\hat{y}_{ui} = f(u, i | \mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u^T \mathbf{q}_i$$



Step 4: NCF Module


- Concatenate outputs from two modules
- Predict scalar values for each input

```
# full neural collaborative filtering module
class NCF(nn.Module):
    def __init__(self, user_num, item_num, output_dim):
        super(NCF, self).__init__()
        self.GMF_layers = GMF(user_num=user_num,
                                item_num=item_num,
                                output_dim=output_dim)
        self.MLP_layers = MLP(user_num=user_num,
                                item_num=item_num,
                                output_dim=output_dim)
        self.predict_layer = nn.Linear(output_dim * 2, 1)
        self.dropout = nn.Dropout()

    def forward(self, user, item):
        # GMF layers
        output_GMF = self.GMF_layers(user, item)

        # MLP layers
        output_MLP = self.MLP_layers(user, item)

        # merge outputs
        concat = torch.cat((output_GMF, output_MLP), -1)
        concat = self.dropout(concat)
        prediction = self.predict_layer(concat).view(-1)
        return prediction
```



Step 4: NCF Module


- Concatenate outputs from two modules
- Predict scalar values for each input

```
# full neural collaborative filtering module
class NCF(nn.Module):
    def __init__(self, user_num, item_num, output_dim):
        super(NCF, self).__init__()
        self.GMF_layers = GMF(user_num=user_num,
                                item_num=item_num,
                                output_dim=output_dim)
        self.MLP_layers = MLP(user_num=user_num,
                                item_num=item_num,
                                output_dim=output_dim)
        self.predict_layer = nn.Linear(output_dim * 2, 1)
        self.dropout = nn.Dropout()

    def forward(self, user, item):
        # GMF layers
        output_GMF = self.GMF_layers(user, item)

        # MLP layers
        output_MLP = self.MLP_layers(user, item)

        # merge outputs
        concat = torch.cat((output_GMF, output_MLP), -1)
        concat = self.dropout(concat)
        prediction = self.predict_layer(concat).view(-1)
        return prediction
```




Step 5: Training Loop

- Initialize model
- Initialize loss and optimizer
- Compute MSE loss and update params

```
model = NCF(user_num, item_num, output_dim)
model = model.to(device)
loss_function = nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

for epoch in range(1, 20):
    # training loop
    start_time = time.time()
    model.train()
    for user, item, label in tqdm(train_loader, total=len(train_loader)):
        user = user.to(device)
        item = item.to(device)
        label = label.float().to(device)
        model.zero_grad()
        prediction = model(user, item)
        loss = loss_function(prediction, label)
        loss.backward()
        optimizer.step()
```



Step 6: Evaluation

- Remember the set model to eval mode
- Compute desired metrics such as MAP

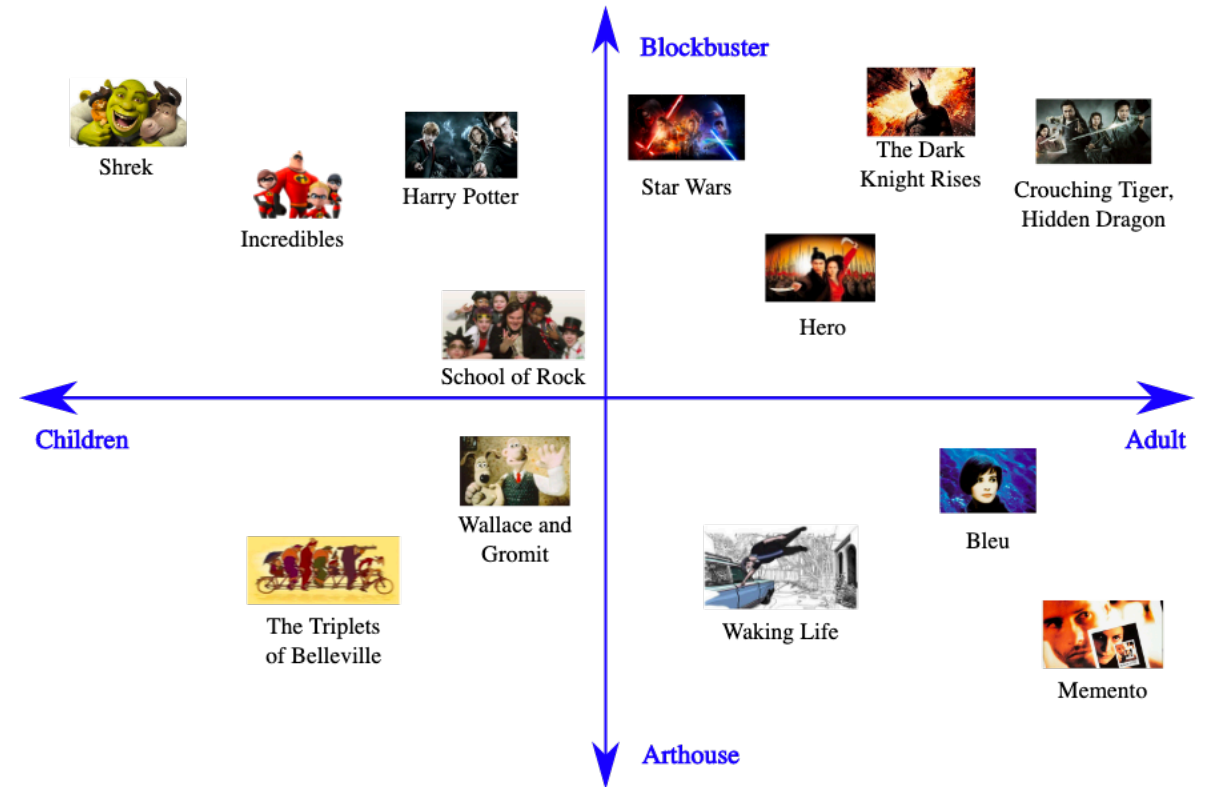
```
# eval loop
test_mae = []
model.eval()
for user, item, label in test_loader:
    user = user.to(device)
    item = item.to(device)
    prediction = model(user, item)
    label = label.float().detach().cpu().numpy()
    prediction = prediction.float().detach().cpu().numpy()
    MAE = mean_absolute_error(y_pred=prediction, y_true=label)
    test_mae.append(MAE)
```

Recommendation

- Using NCF, we trained a model to directly predict ratings of (user, item) pairs
- Predict ratings of unobserved items
 - Can be used as a ranking function
 - High rating can be potentially similar items

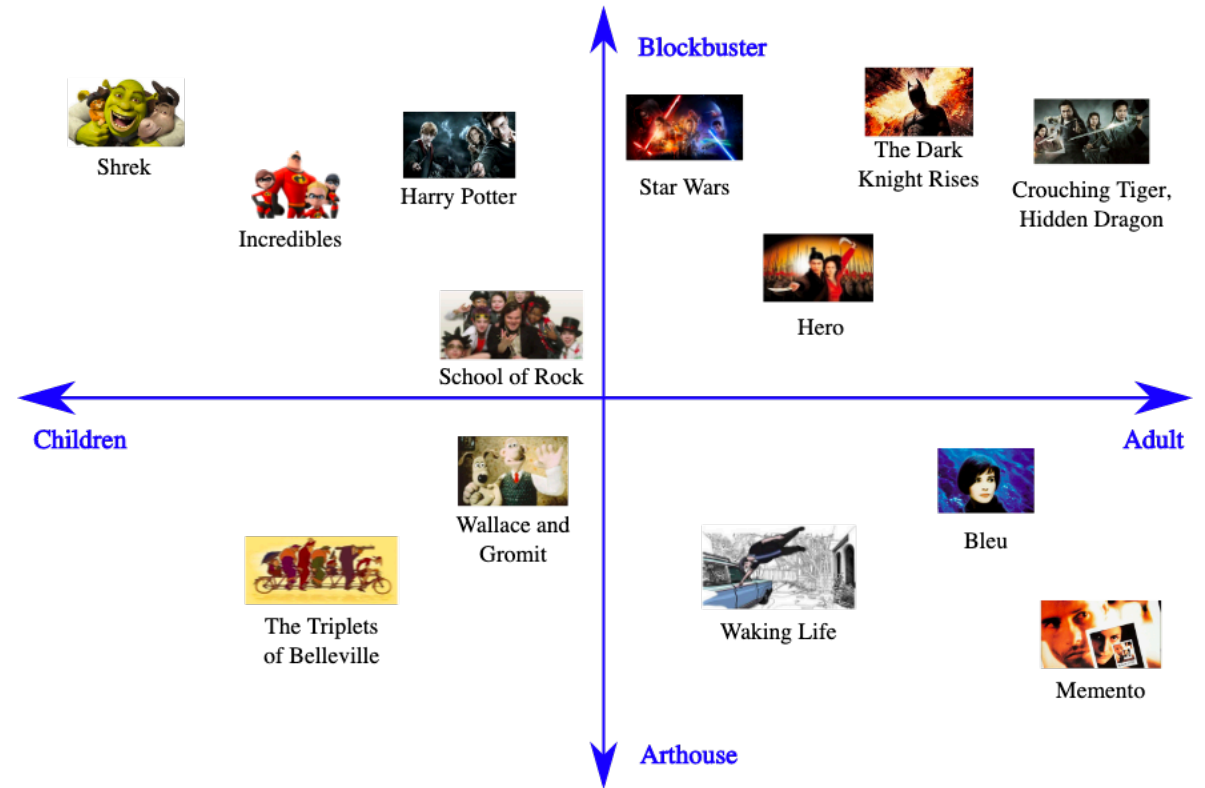
Item Similarity

- Projected one-hot item vectors into dense embeddings
- Contain notion of similarity
- Dimensionality reduction algorithms can be applied for visualization



Obtaining Weights of PyTorch Layer

- `model.parameters()`
- `model.layer.weight`
- Convert to numpy array and apply T-SNE or PCA



Coding Exercise

- https://github.com/boslbi92/pytorch_tutorial
- Implement class simpleCF inside **practice.py**
- Solution is available at **solution.py**